

# Stable Diffusion and Paper Reading

DreamBooth: Fine Tuning Text-to-Image Diffusion Models for Subject-Driven Generation

Junzhe Yi

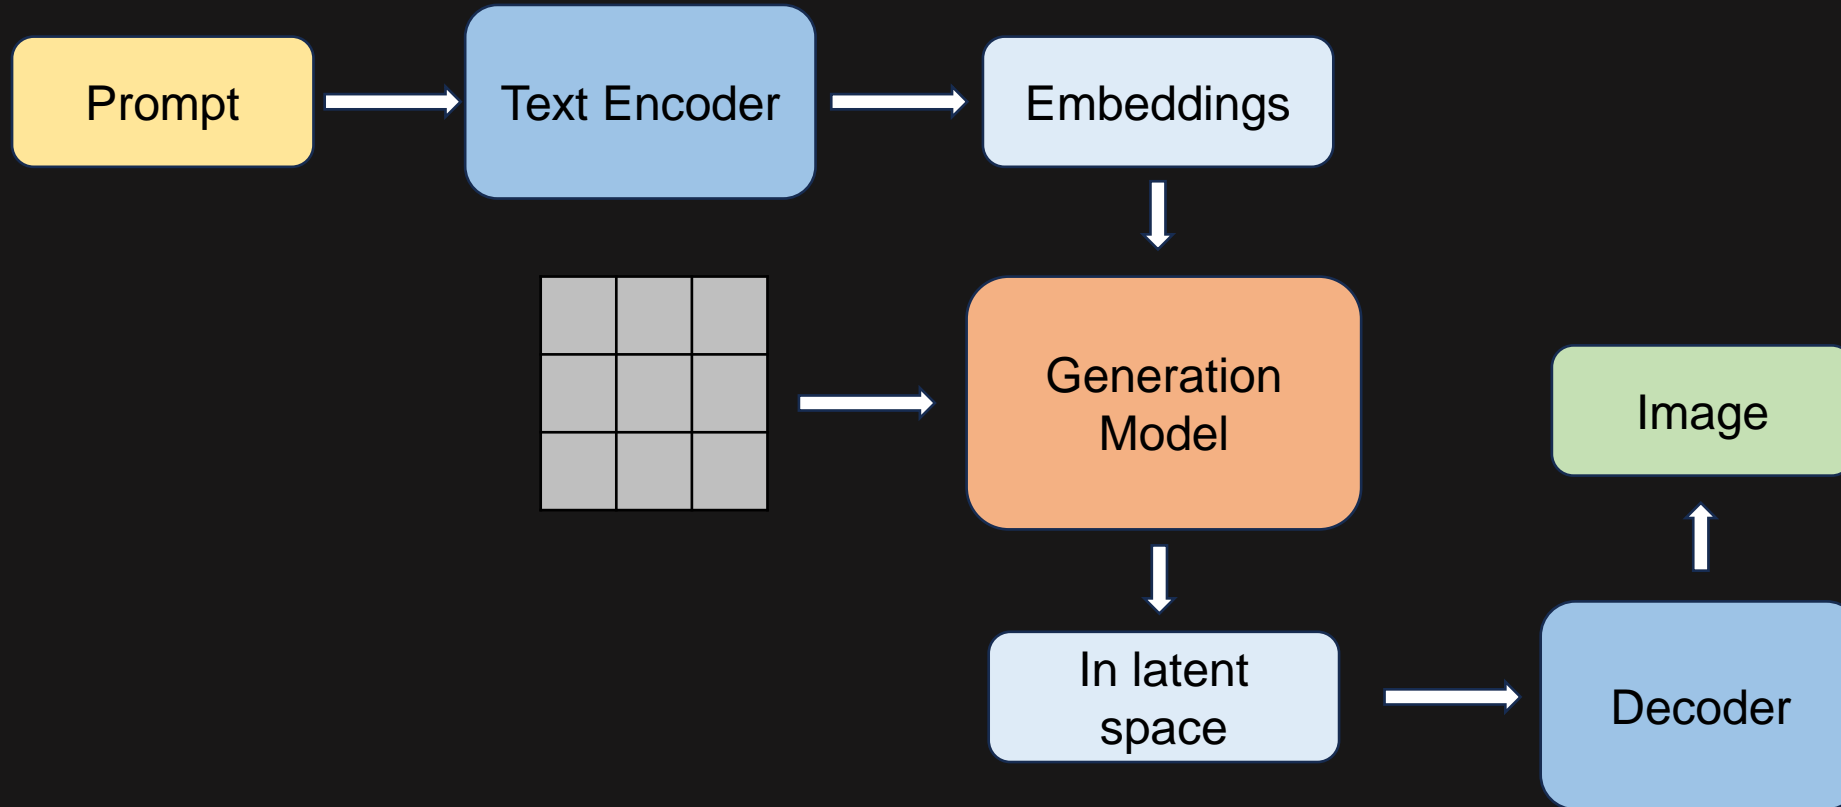

HuazhongUST

April 27, 2024

# Summary

1. Framework of stable diffusion
2. Understanding how diffusion model works
3. Fine tuning: Dreambooth
4. Methodology Overview
5. Reference

# Framework of Diffusion Models



# Framework of Stable Diffusion

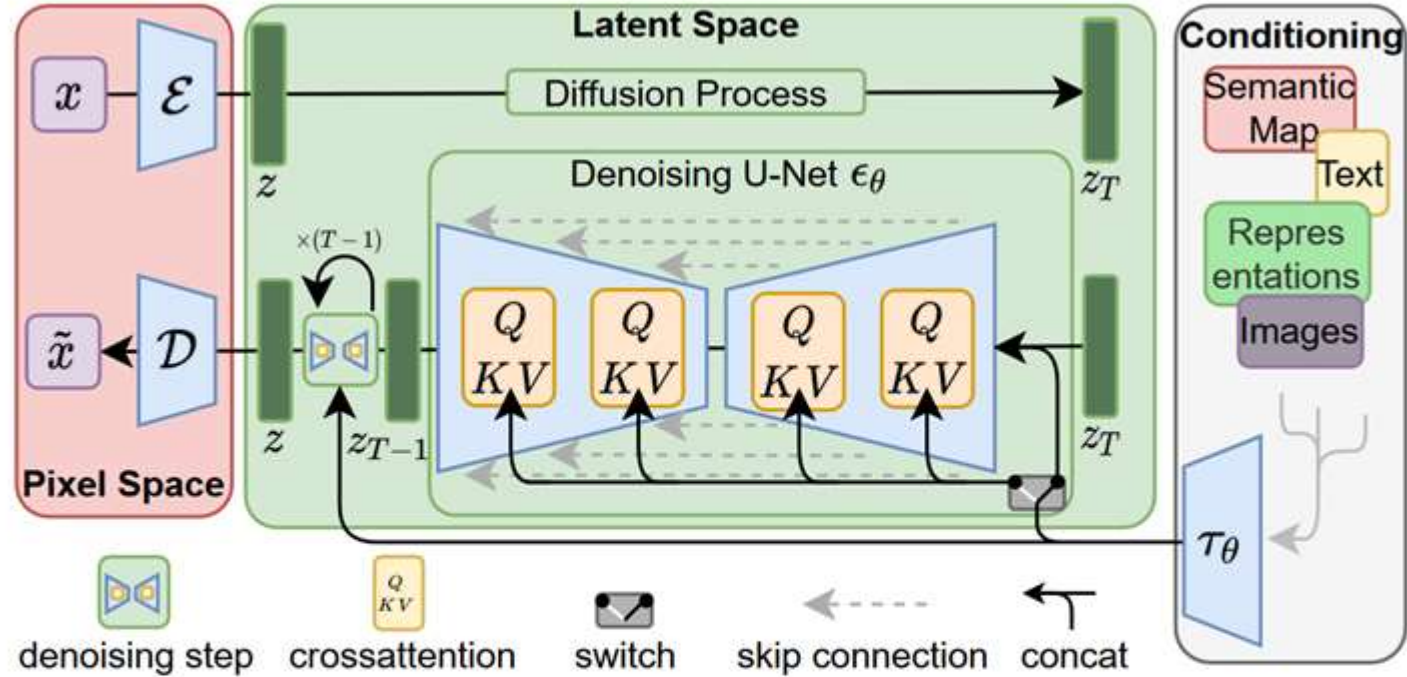
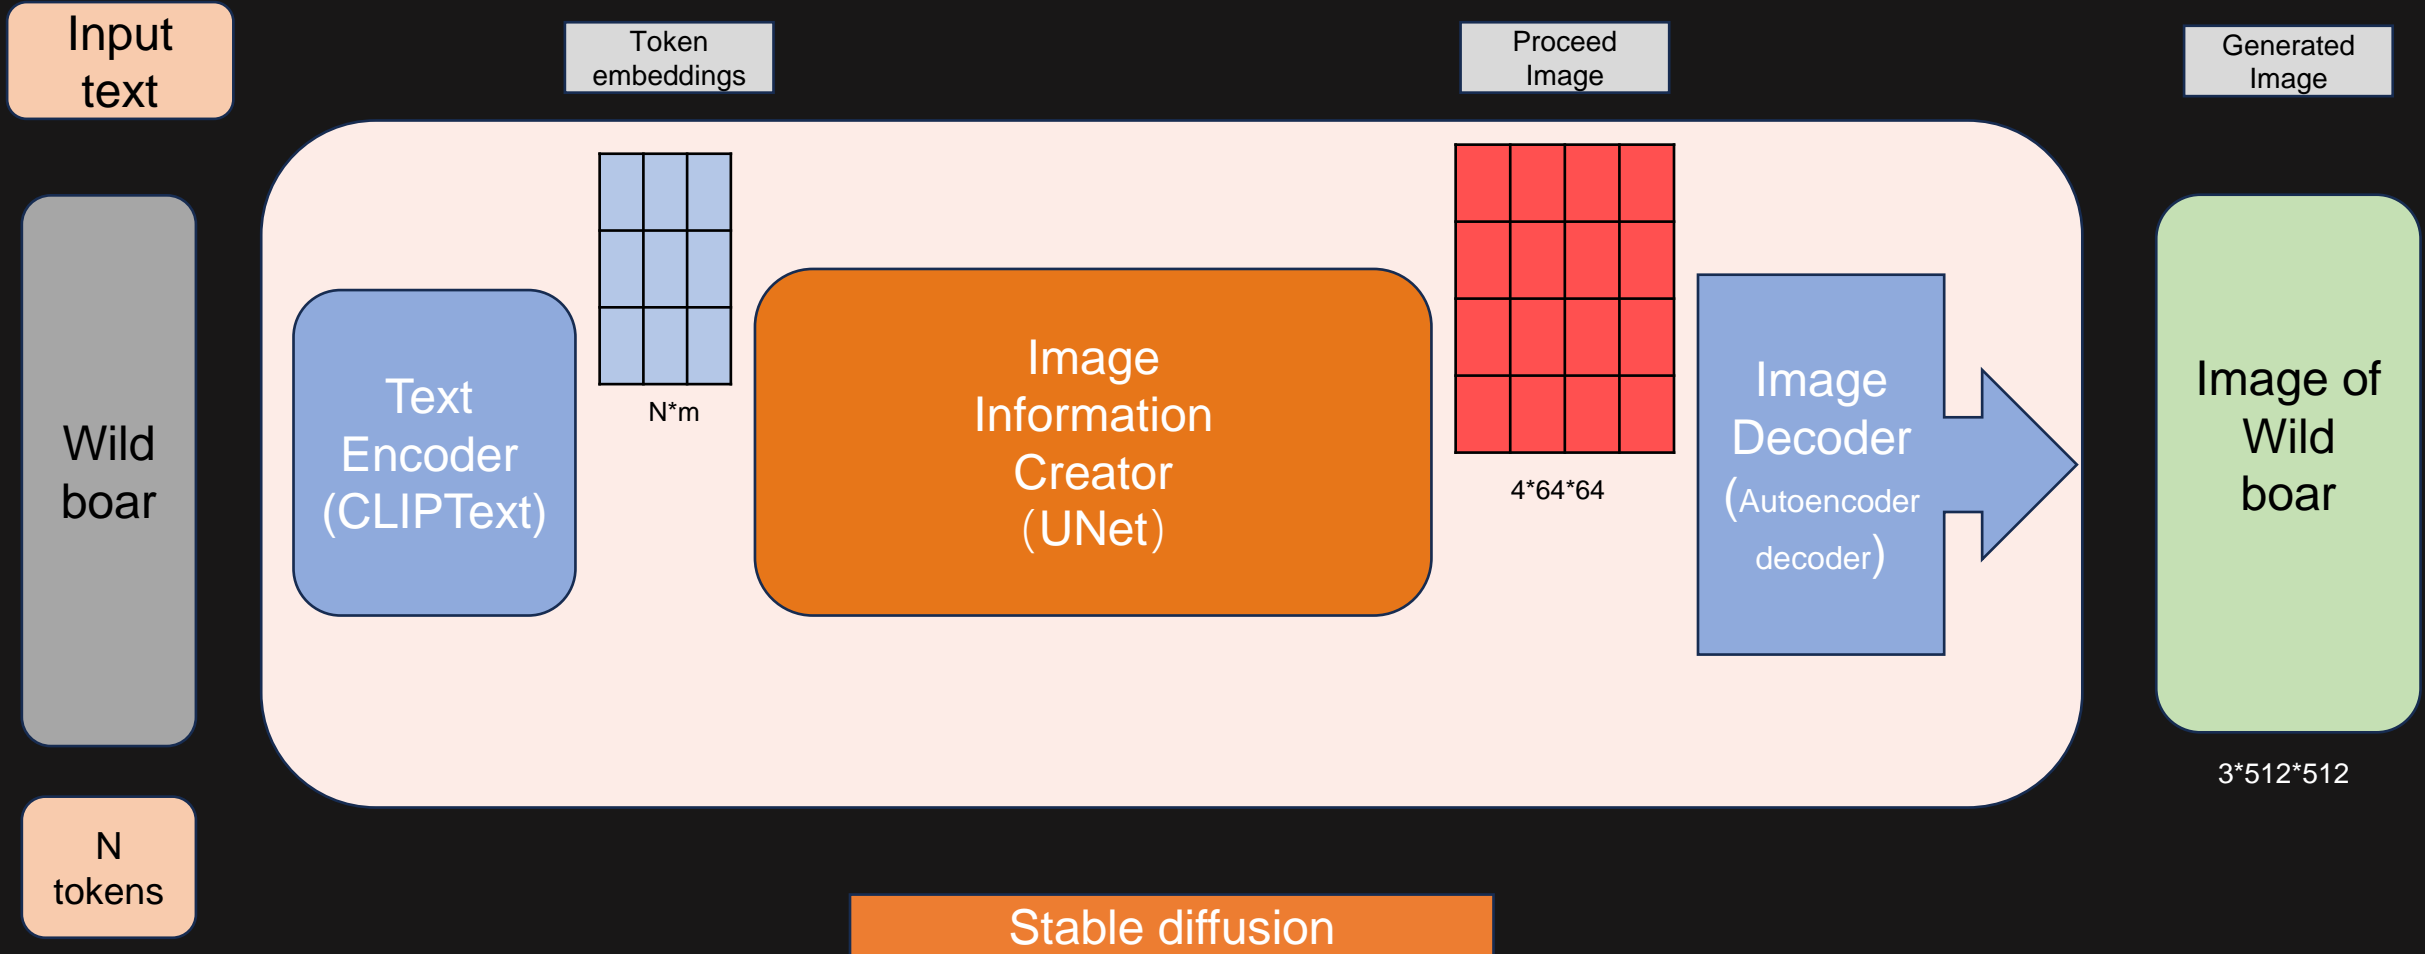


Figure 3. We condition LDMs either via concatenation or by a more general cross-attention mechanism. See Sec. 3.3

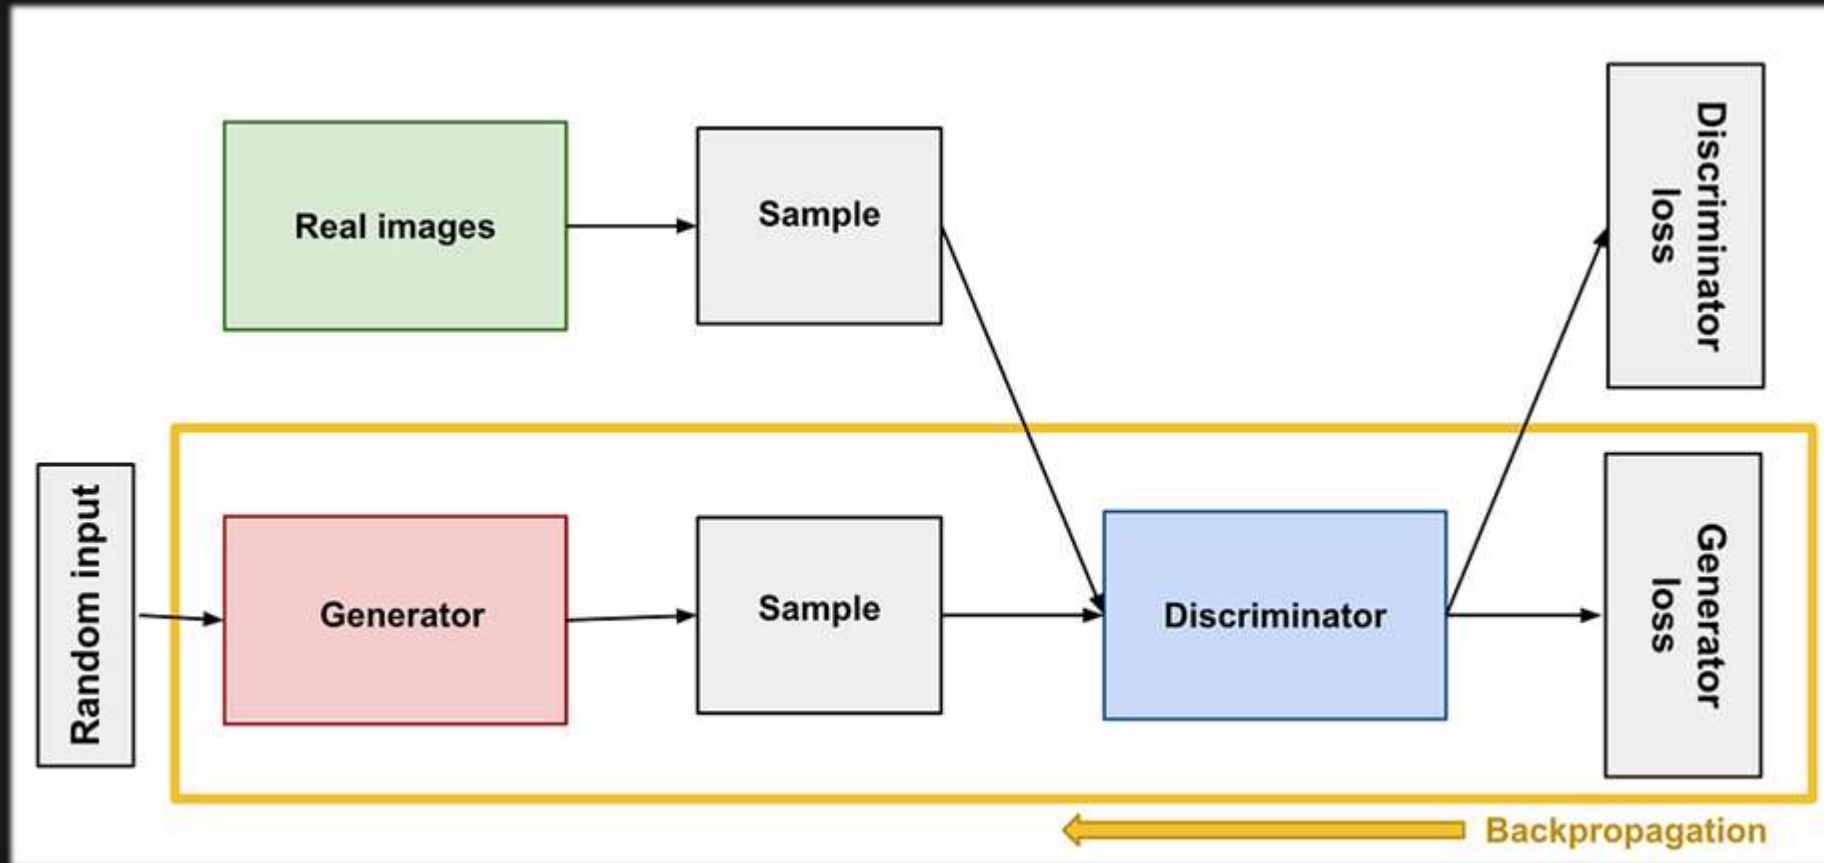
*High-Resolution Image Synthesis With Latent Diffusion Models* (2022 CVPR)  
<https://arxiv.org/abs/2112.10752>

# Framework of Stable Diffusion



# How Image Creator Work?

Generative adversarial network ( GAN )



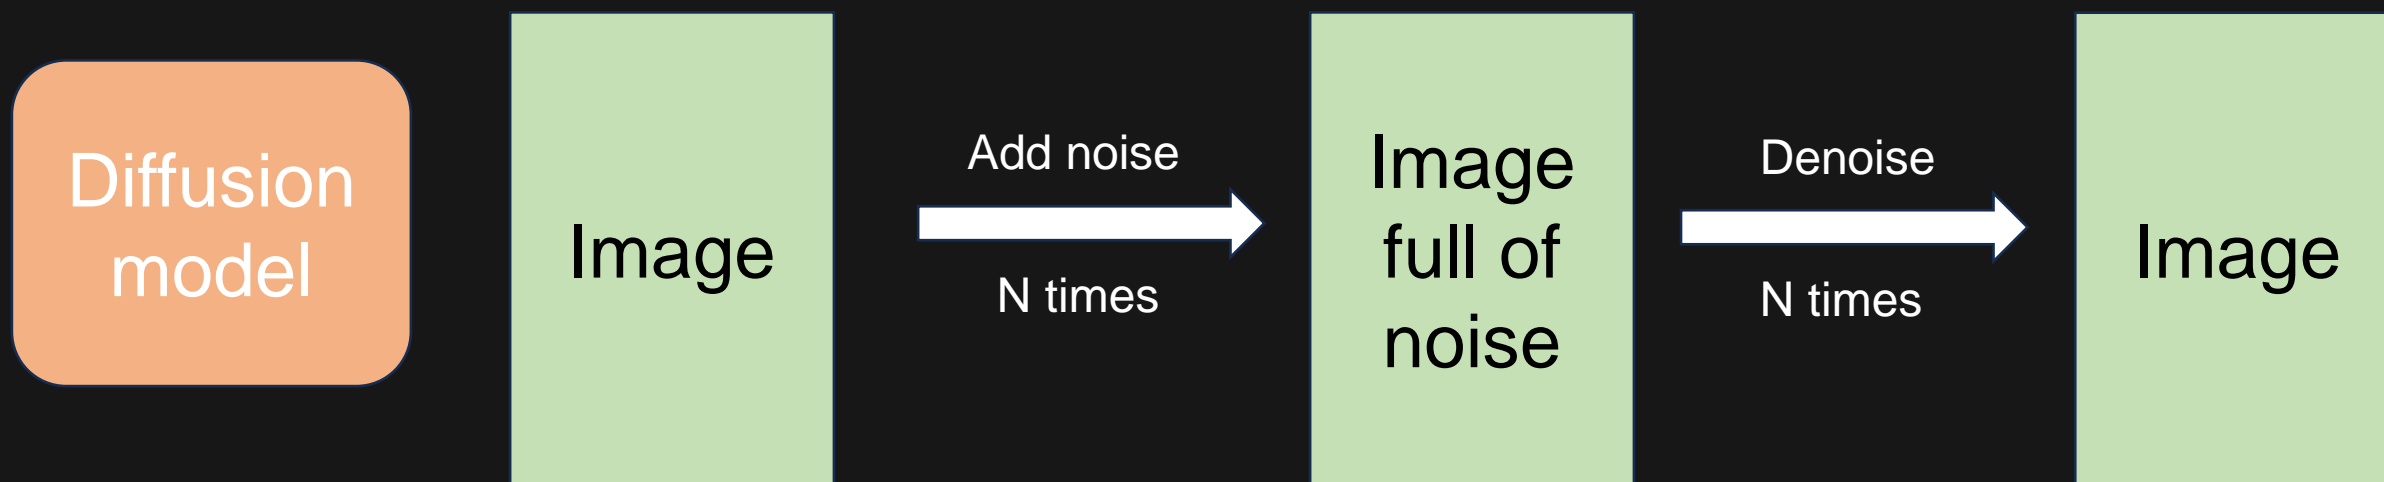
*Generative Adversarial Nets*  
<https://arxiv.org/abs/1406.2661>

# DDPM

Diffusion model

DDPM (NeurIPS2020)

They work by gradually adding Gaussian noise to the original data in the forward diffusion process and then learning to remove the noise in the reverse diffusion process.

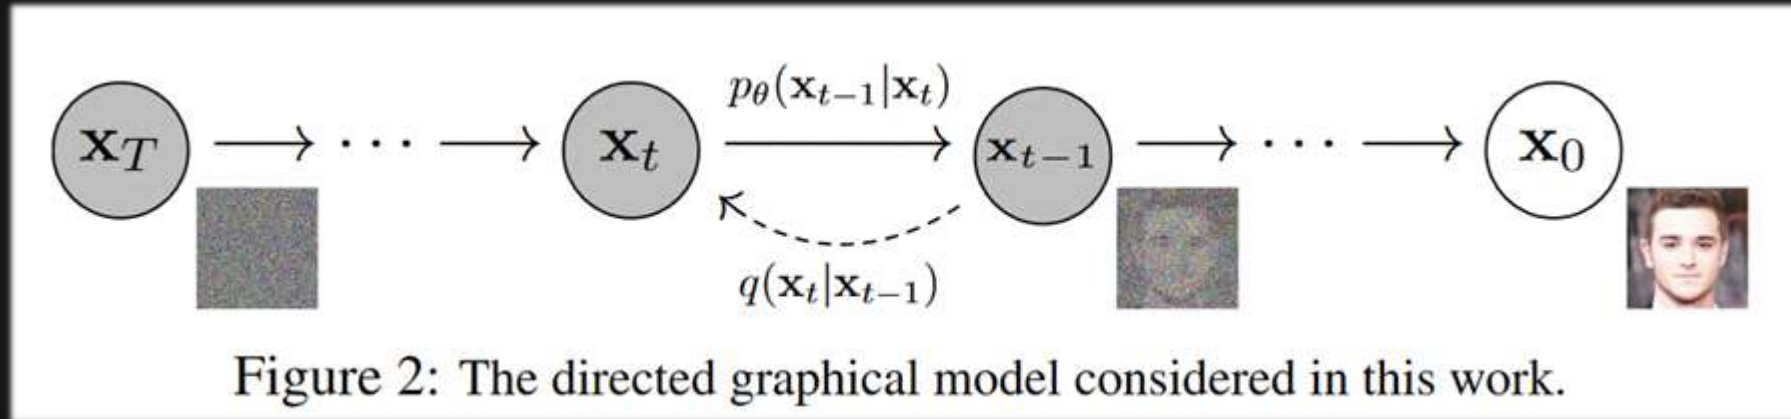


*Denoising Diffusion Probabilistic Models*  
<https://arxiv.org/abs/2006.11239>

# Two Step Process

A denoising diffusion modeling is a two step process:

- Forward diffusion process — The forward diffusion process is the Markov chain of diffusion steps in which we slowly and randomly add noise to the original data.
- Reverse diffusion process — The reverse diffusion process tries to reverse the diffusion process to generate original data from the noise.





# Algorithm

---

**Algorithm 1** Training

---

```
1: repeat  
2:    $\mathbf{x}_0 \sim q(\mathbf{x}_0)$   
3:    $t \sim \text{Uniform}(\{1, \dots, T\})$   
4:    $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
5:   Take gradient descent step on  
        $\nabla_{\theta} \|\boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t)\|^2$   
6: until converged
```

---

---

**Algorithm 2** Sampling

---

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$   
2: for  $t = T, \dots, 1$  do  
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$   
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$   
5: end for  
6: return  $\mathbf{x}_0$ 
```

---

# Understanding Diffusion in Mathematical Detail

Markov chain:  $x_t = \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_{t-1}$        $x_{t-1} = \sqrt{\alpha_{t-1}}x_{t-2} + \sqrt{1 - \alpha_{t-1}}\epsilon_{t-2}$

We can arrive  $x_t = \sqrt{\alpha_t}(\sqrt{\alpha_{t-1}}x_{t-2} + \sqrt{1 - \alpha_{t-1}}\epsilon_{t-2}) + \sqrt{1 - \alpha_t}\epsilon_{t-1}$

That is  $x_t = \sqrt{\alpha_t\alpha_{t-1}}x_{t-2} + \sqrt{1 - \alpha_t\alpha_{t-1}}\bar{\epsilon}_{t-2}$

Conclusion:  $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon_t$

Forward diffusion process finished.

Using Bayes' rule, we have:  $q(x_{t-1}|x_t, x_0) = \frac{q(x_t|x_{t-1}, x_0)q(x_{t-1}|x_0)}{q(x_t|x_0)}$

Given the condition of knowing  $x_t$ , find  $x_{t-1}$ , This probability follows a normal distribution

$$\tilde{\mu}_t(x_t, x_0) = \frac{\sqrt{\alpha_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}x_t + \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t}x_0$$

# Understanding Diffusion in Mathematical detail

The mean of this normal distribution is correlated with  $X_t$  and  $X_0$ . This is going to be difficult now, because there is a **unknown  $X_0$** . However If we go back to the previous page, we will find that  $x_0$  has already been calculated in the forward process.

so we can estimate  $X_0$  as: 
$$x_0 = \frac{1}{\sqrt{\bar{\alpha}_t}} (x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_t)$$

The mean: 
$$\tilde{\mu}_t = \frac{1}{\sqrt{a_t}} \left( x_t - \frac{\beta_t}{\sqrt{1 - \bar{a}_t}} \epsilon_t \right)$$

In this equation, **there is only one unknown**, which is:  $\epsilon_t$

**the next step: to estimate the noise!**

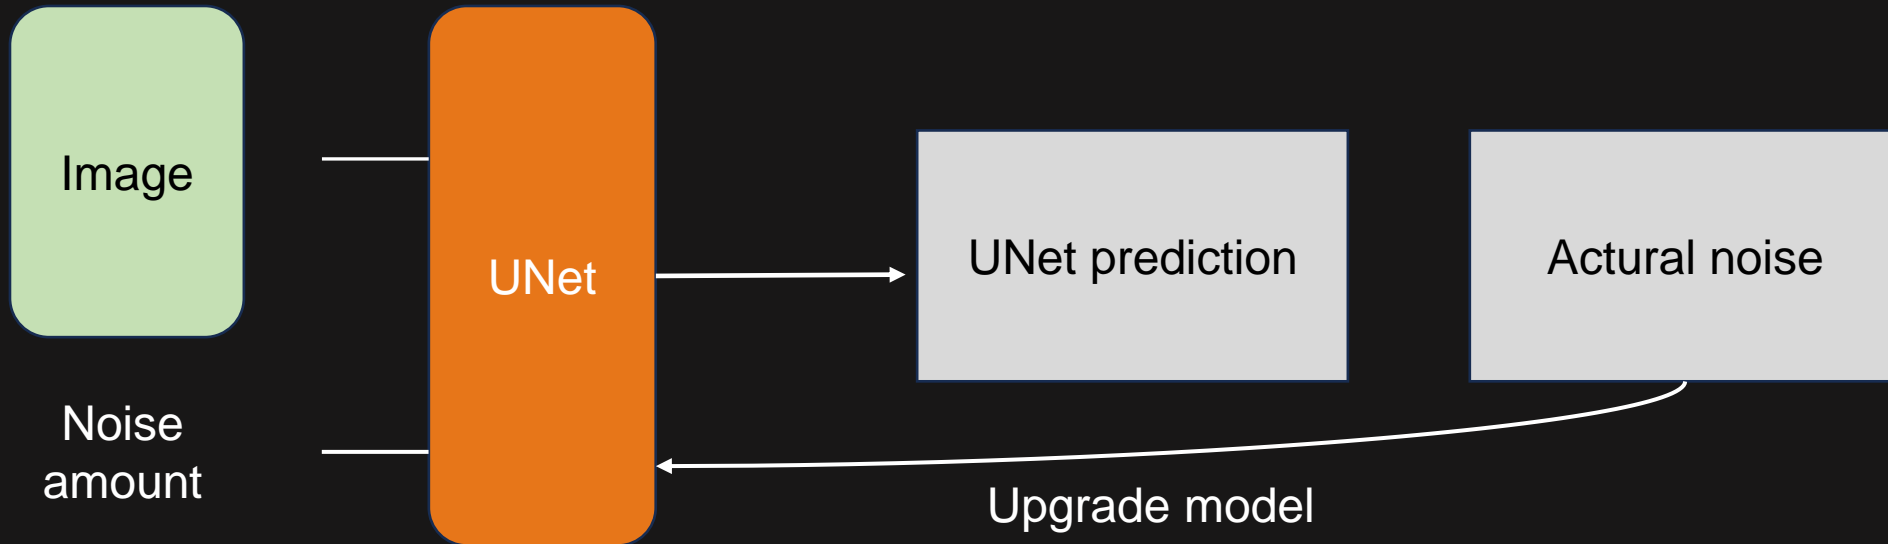

Usually we use 'U-Net + transformer'. Trained on such a massive dataset, the powerful noise predictor U-Net has the "capability" to iteratively denoise a noisy image during the reverse process of diffusion, transforming it into a perfect image.

**Why not have a look into U-Net?**

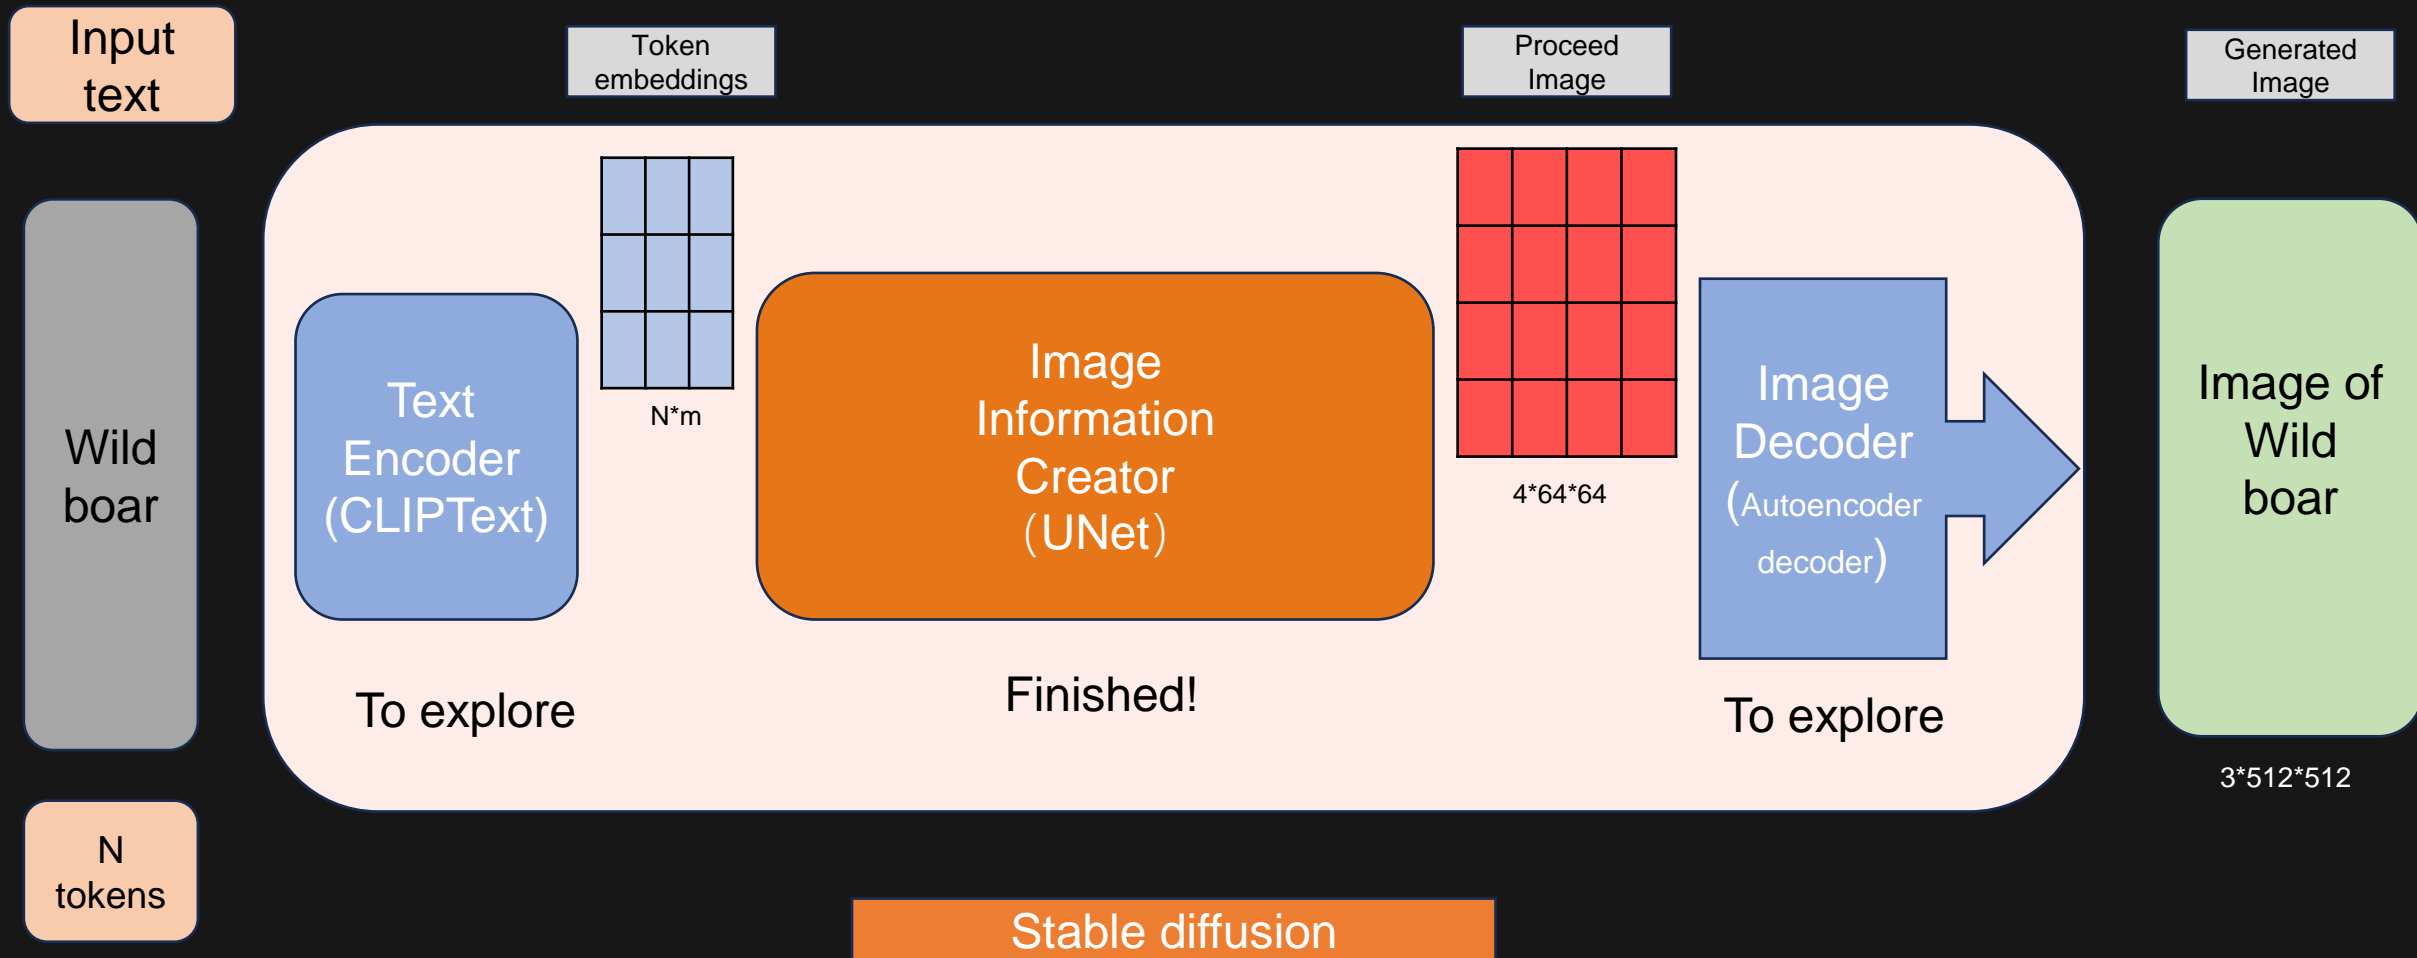
# UNet Train Steps

U-Net train step:

supervised learning!



# Take a Look Back



# AutoEncoder Decoder

To understand AutoEncoder Decoder, you need to know “latent space” as prior knowledge. if you don't, please refer to the LDM, the paper's link is given in Page 5 of this Powerpoint.

In order to discover the latent connections and patterns between images and reduce computational complexity, the operation of Stable Diffusion is not conducted directly on the pixel dimensions of the images themselves but rather in the compressed version of the images, known as the "latent space." This compression and decompression process is achieved through an Autoencoder Decoder.

# AutoEncoder Decoder

## Departure to Latent Space

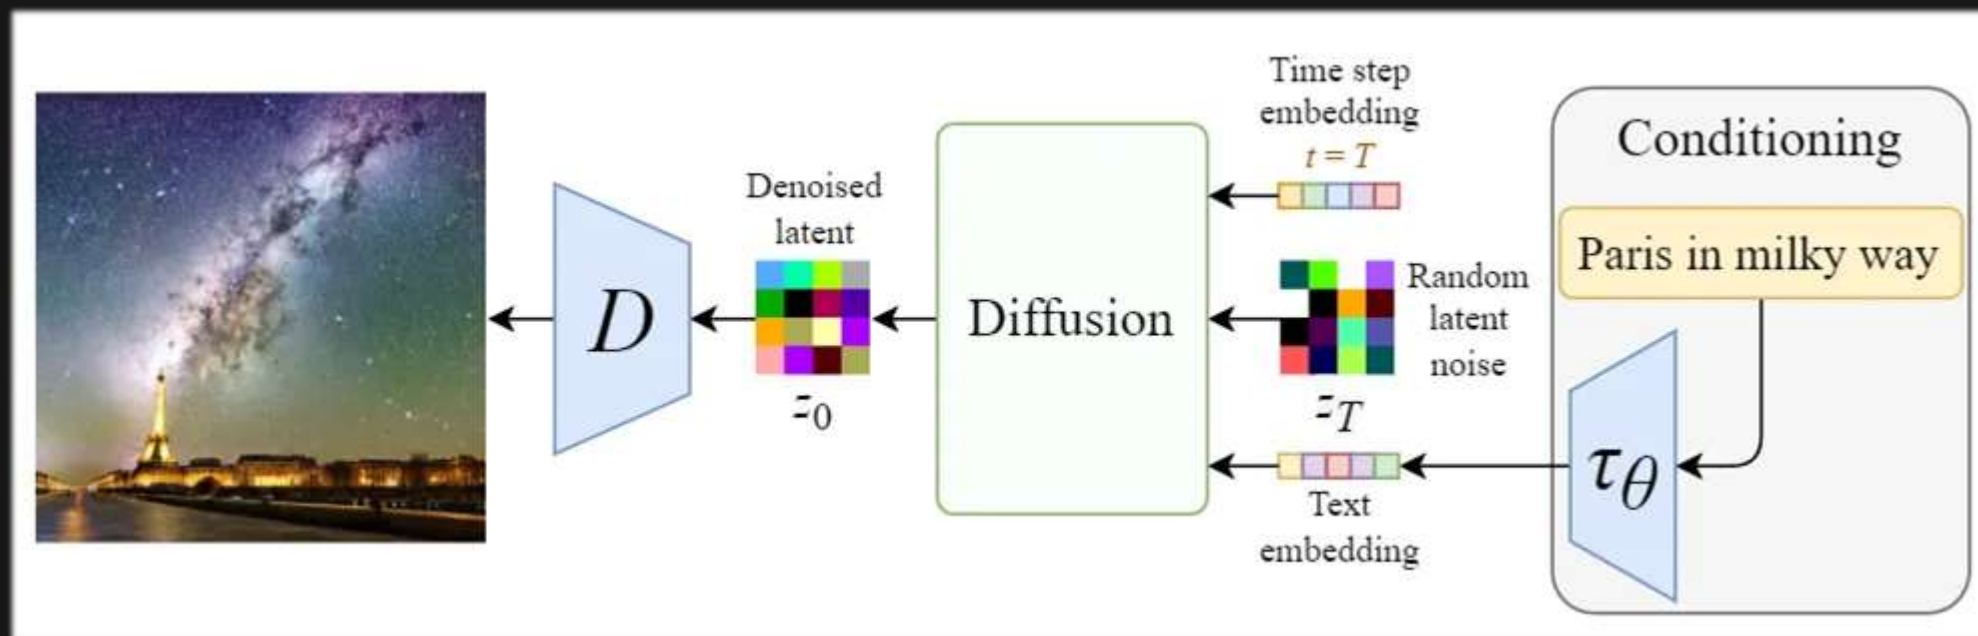
- By using the trained encoder  $E$ , we can encode the full-sized image into lower dimensional latent data (compressed data).
- By using the trained decoder  $D$ , we can decode the latent data back into an image.

## Latent Diffusion

- After encoding the images into latent data, the forward and reverse diffusion processes will be done in the latent space.

# Example

An example





# CLIPtext (Contrastive Language-Image Pre-Training)

At this point, let us introduce the most powerful module of Stable Diffusion.

If you don't have CLIPtext, you won't be able to use textual prompts to control the semantic content of the output image. In that case, Diffusion will generate images randomly.

CLIPtext is a Text Encoder, represented by the deep blue module in the diagram. It is a special type of Transformer-based natural language model. It takes the input text prompts and generates the Token embeddings matrix.

You can understand it as a conditioning mechanism.

# Conditioning of Stable Diffusion

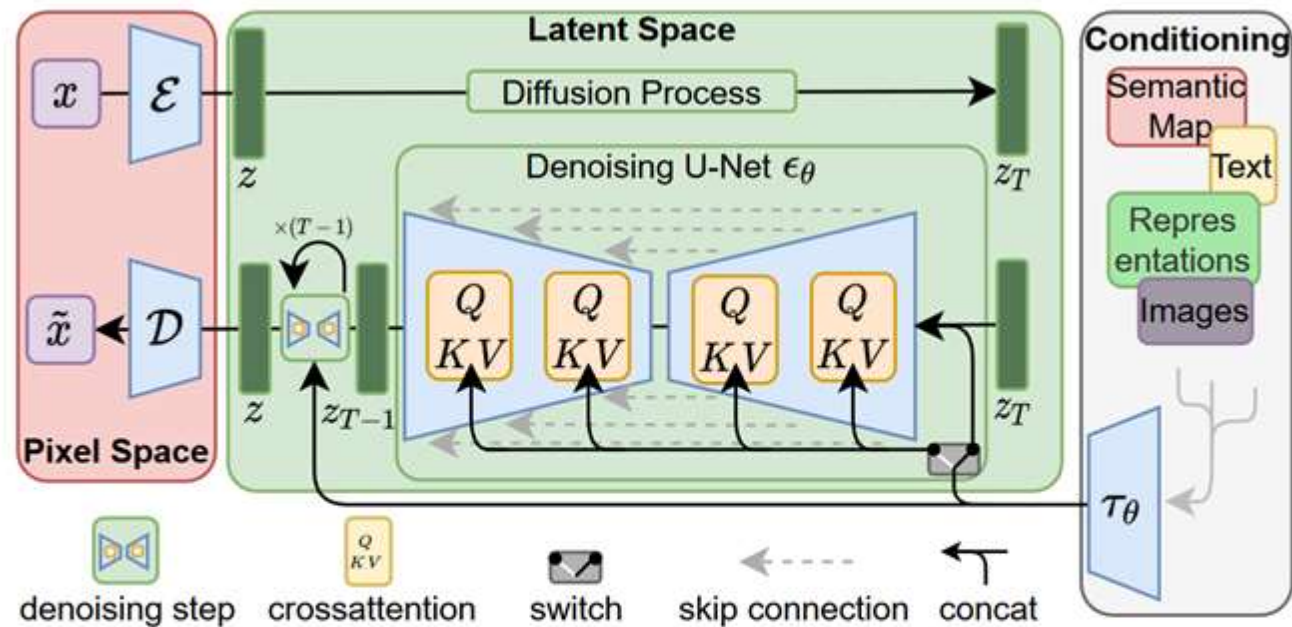


Figure 3. We condition LDMs either via concatenation or by a more general cross-attention mechanism. See Sec. 3.3

This is done by modifying the inner diffusion model to accept conditioning inputs.

# Conditioning of Stable Diffusion

The inner diffusion model is turned into a conditional image generator by augmenting its denoising U-Net with the cross-attention mechanism.

The switch in the above diagram is used to control between different types of inputs:

- For text inputs, they are first converted into embeddings (vectors) using a language model  $\tau_\theta$  (e.g. BERT, CLIP), and then they are mapped into the U-Net via the (multi-head) Attention(Q, K, V) layer.
- For other spatially aligned inputs (e.g. semantic maps, images, inpainting), the conditioning can be done using concatenation.

# DreamBooth

The current text-to-image generation models are capable of generating high-quality images based on given prompts. However, these models are unable to mimic the appearance of objects in a given reference image and generate new images in different scenes.

Until DreamBooth was introduced.

## **DreamBooth: Fine Tuning Text-to-Image Diffusion Models for Subject-Driven Generation**

Nataniel Ruiz<sup>\*,1,2</sup>

Yuanzhen Li<sup>1</sup>

Varun Jampani<sup>1</sup>

Yael Pritch<sup>1</sup>

Michael Rubinstein<sup>1</sup>

Kfir Aberman<sup>1</sup>

<sup>1</sup> Google Research    <sup>2</sup> Boston University

# Identified Gaps


- Most text-to-image generation models are unable to perform few-shot learning
  - There is no efficient method to “inject” subjects into the model training
    - Overfitting
    - Language Drift
    - Best ex.: GAN reproducing the same face given ~100 images[1]

[1] Karras, Tero, Samuli Laine, and Timo Aila. "A style-based generator architecture for generative adversarial networks." Proceedings of the IEEE/CVF conference on computer vision and pattern recognition. 2019

# Contributions

- New efficient few-shot fine-tuning technique that preserves semantic class knowledge
  - Tuning requires only 3~5 examples
- Created new problem to explore diffusion model capabilities-subject-driven generation
  - Goal: maintain fidelity in new contexts

# Improvements over Existing Work



Input Images



Image-guided, DALL-E2



Text-guided, Imagen



Ours



# Subject-Driven Generation (Personalization)



Input images



in the Acropolis



swimming




in a doghouse



sleeping



in a bucket

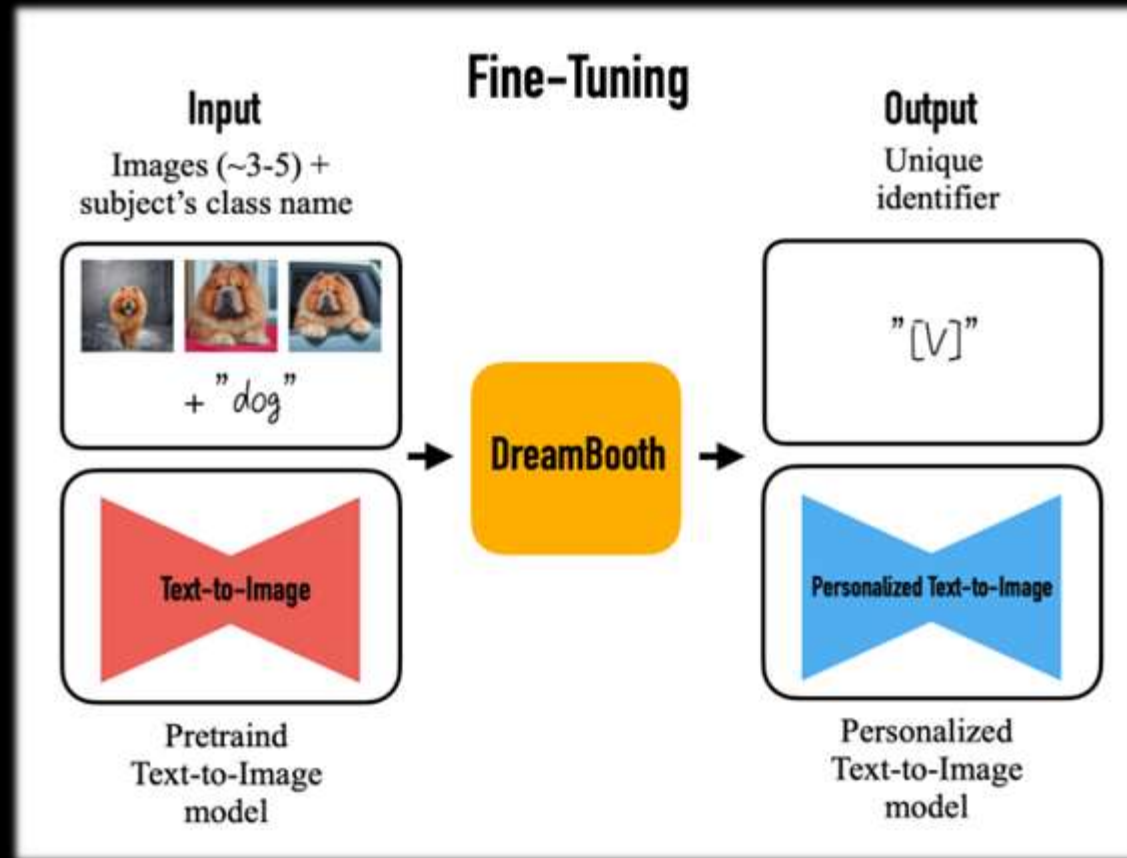


getting a haircut



# Methodology Overview

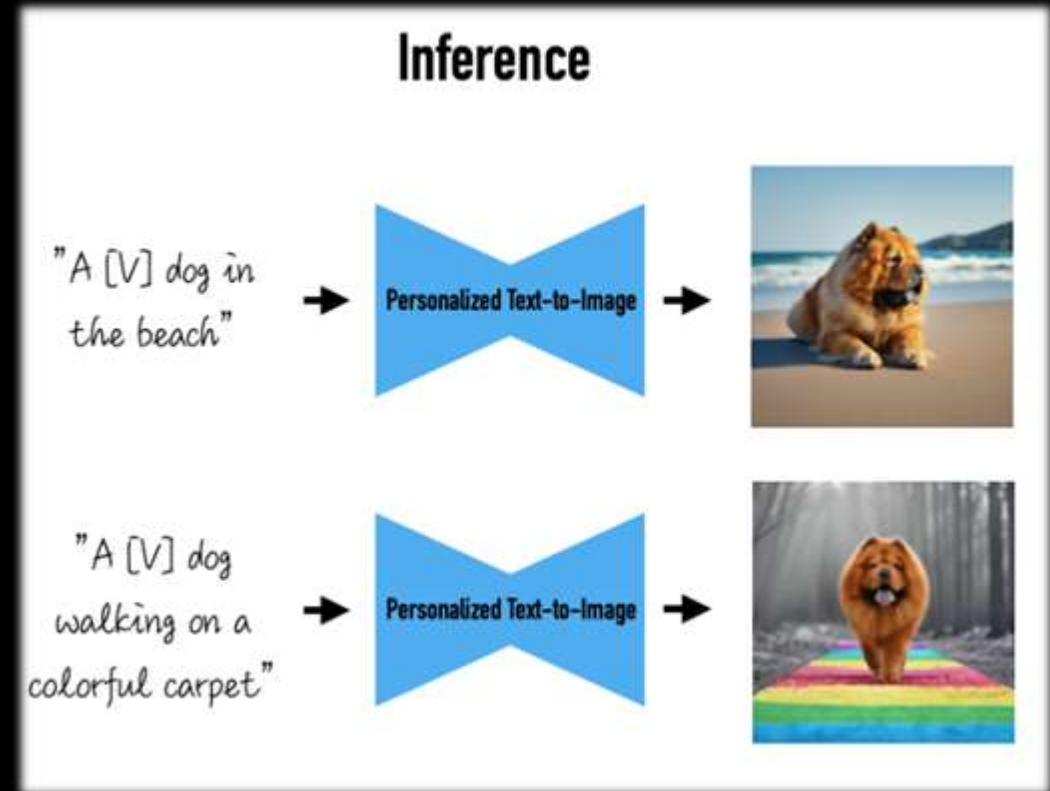
- Few-shot text-guided diffusion
  - Fine-tune existing diffusion models



# Rare-Token Identifier

? There is no efficient method to “inject” subjects into the model training

- How to refer to a new subject?
  - Image is straightforward
- Diffusion model already has a ‘vocabulary’
  - Goal: Implant the subject into this vocabulary

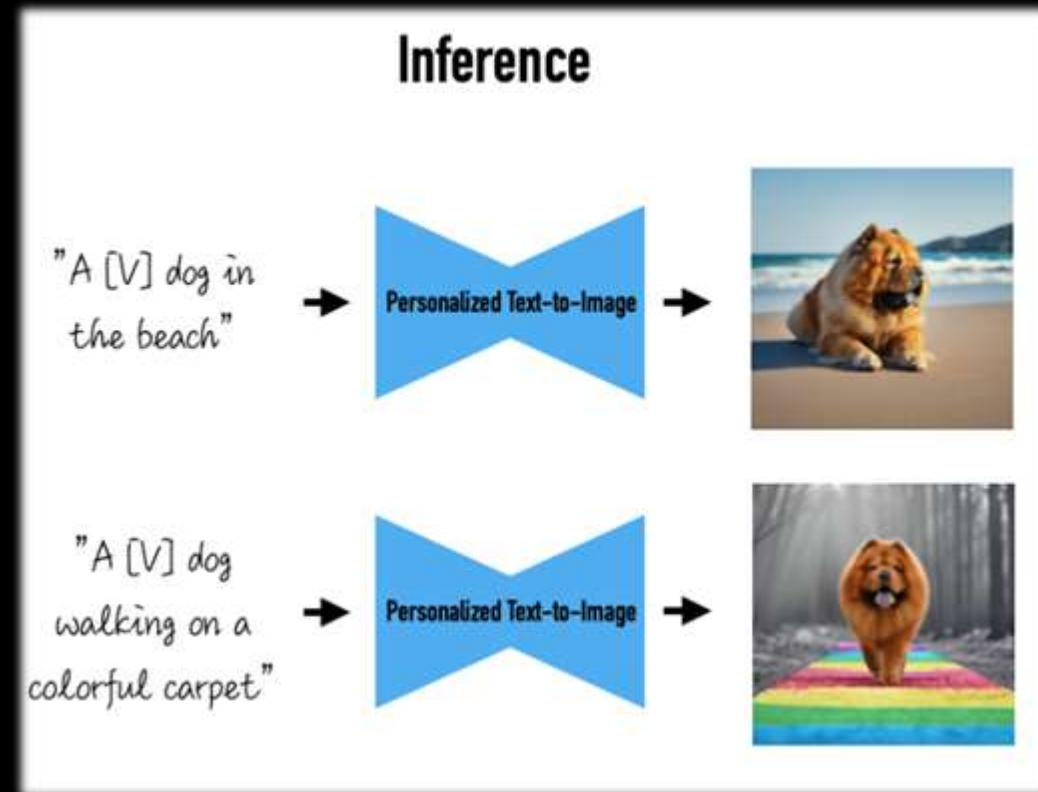


# Rare-Token Identifier: Native Approaches

- Use English words to describe the new class
  - English lexicon has strong priors 'unique' or 'special'
- Generate a string of random characters
  - Nonsense identifier leads to literal artifacts e.g. "xxy5syt00"

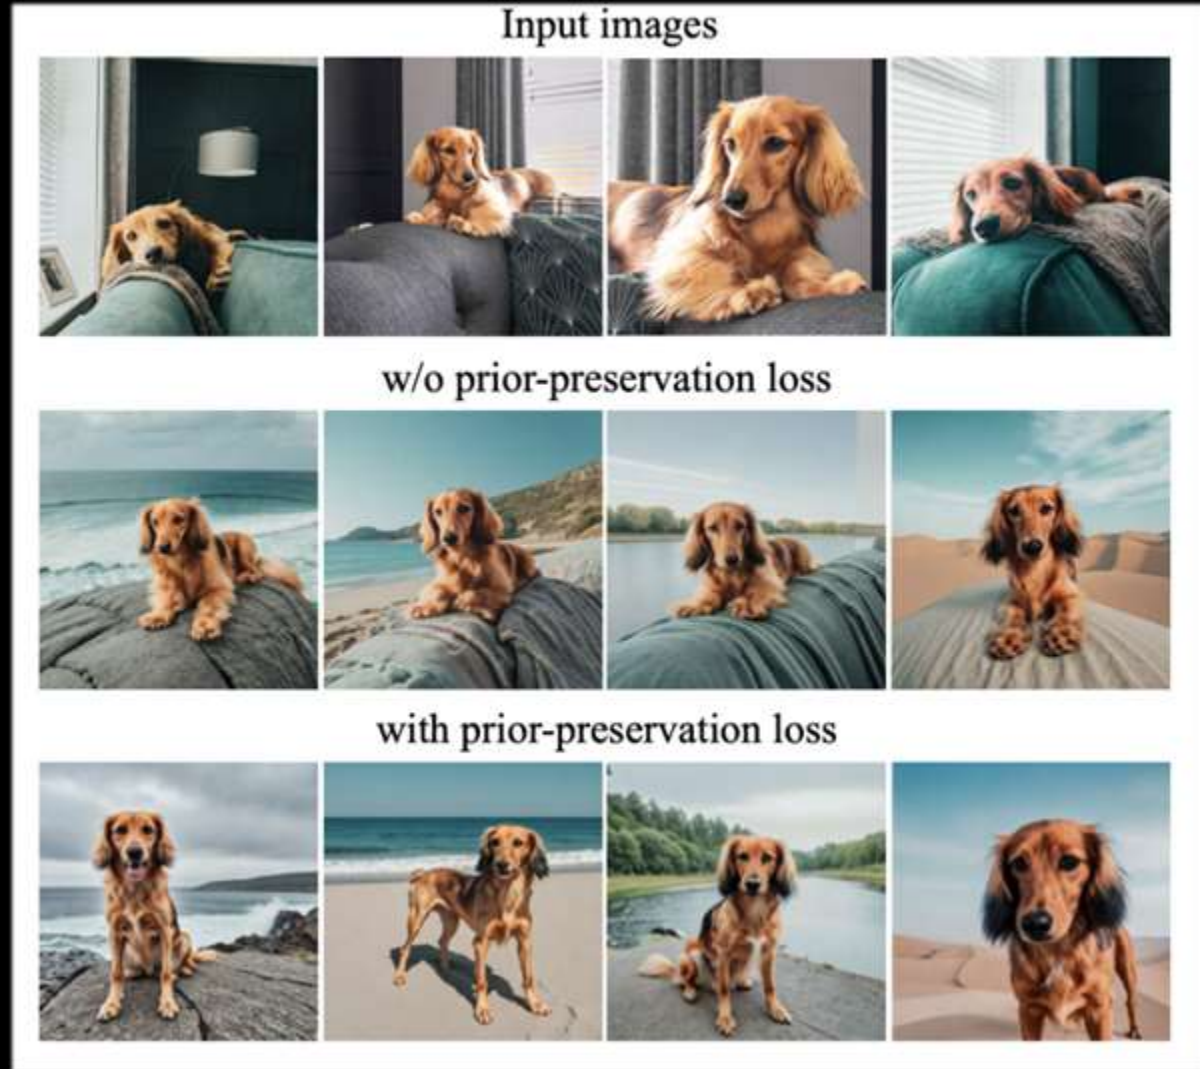
# Rare-Token Identifier: Solution

1. Perform a rare-token lookup in the vocabulary
2. Invert the rare token, resulting in the plain text. (1~3 characters)
3. Use the plain text as the unique identifier



# Autogenous Class-Specific Prior Preservation Loss

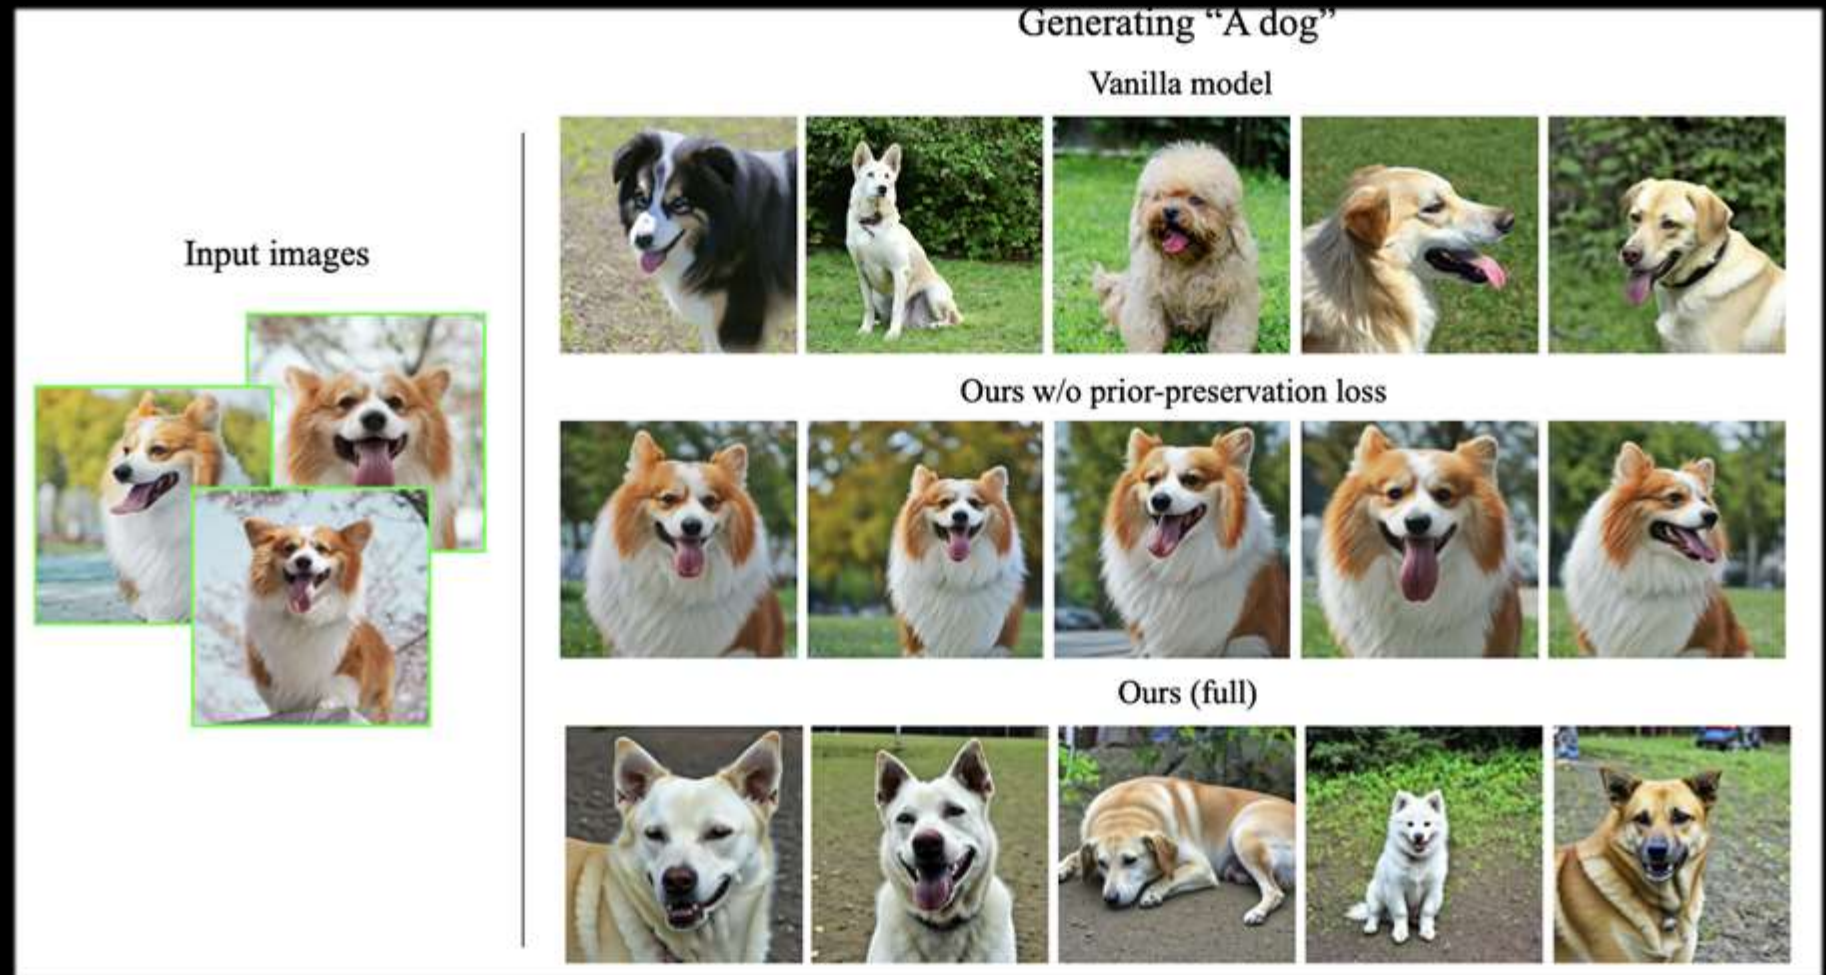
## 1. Overfitting



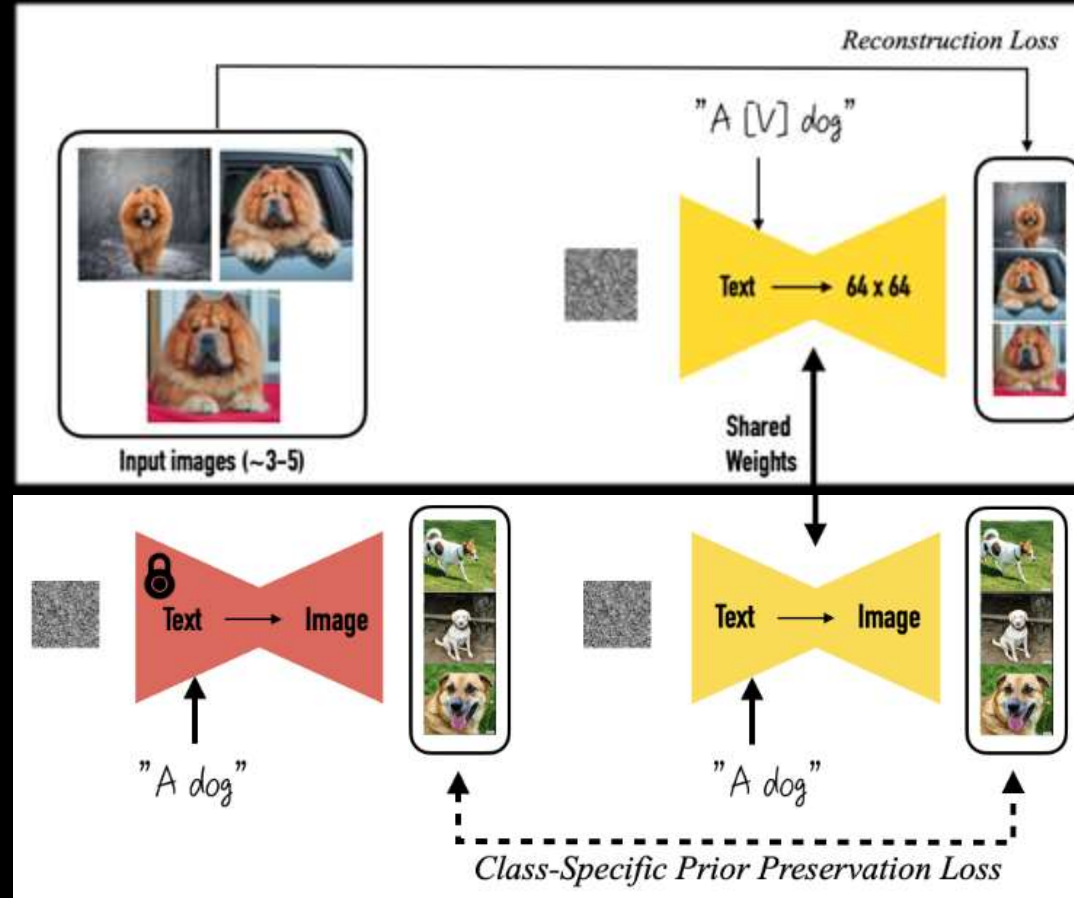


# Autogenous Class-Specific Prior Preservation Loss

## 2. Language Drift



# Autogenous Class-Specific Prior Preservation Loss



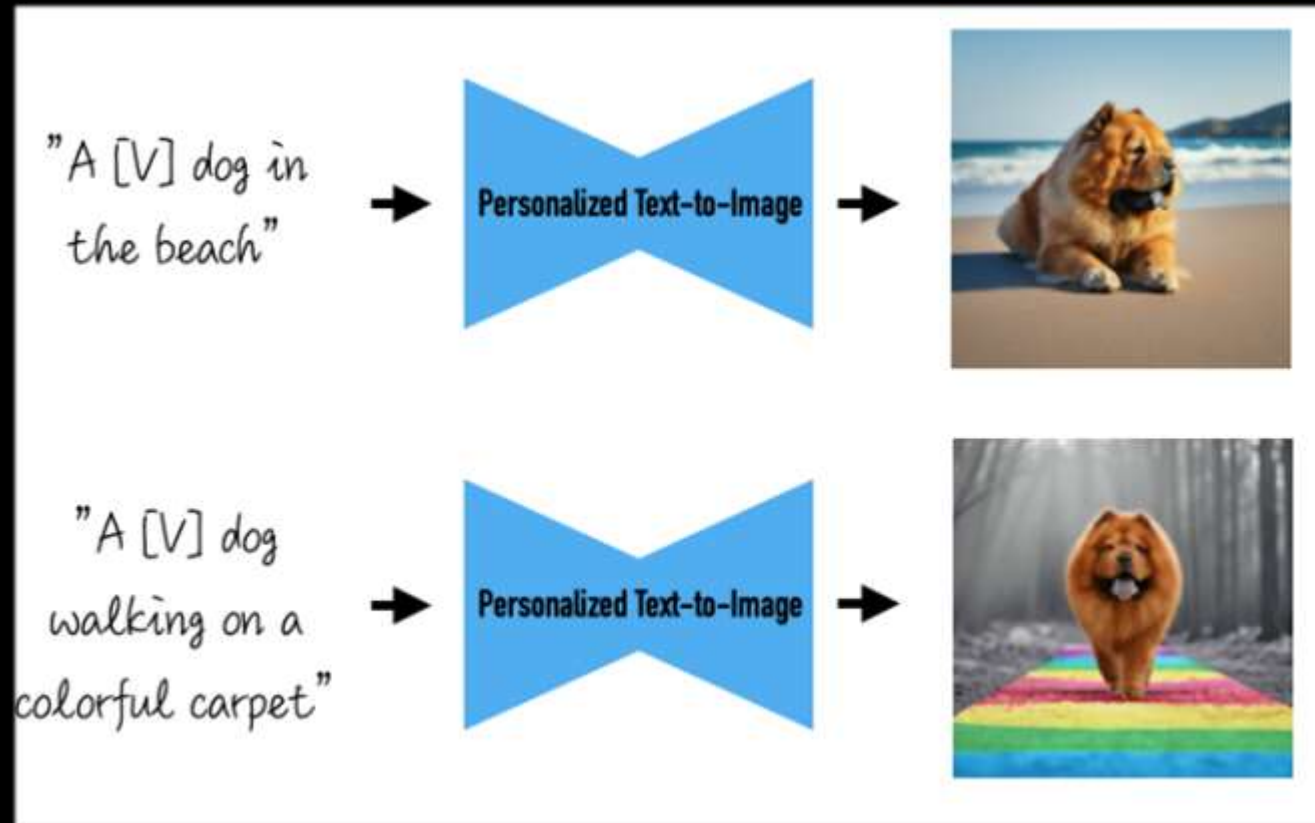
Standard Diffusion Loss

$$\mathbb{E}_{\mathbf{x}, \mathbf{c}, \epsilon, t} [w_t \|\hat{\mathbf{x}}_\theta(\alpha_t \mathbf{x} + \sigma_t \epsilon, \mathbf{c}) - \mathbf{x}\|_2^2]$$

Autogenous Class-Specific Prior Preservation Loss

$$+\lambda w_{t'} \|\hat{\mathbf{x}}_\theta(\alpha_{t'} \mathbf{x}_{\text{pr}} + \sigma_{t'} \epsilon', \mathbf{c}_{\text{pr}}) - \mathbf{x}_{\text{pr}}\|_2^2]$$

# Results





# Results: Recontextualization



# Results: Art Renditions

Input images



Vincent Van Gogh



Michelangelo



Rembrandt



Johannes Vermeer



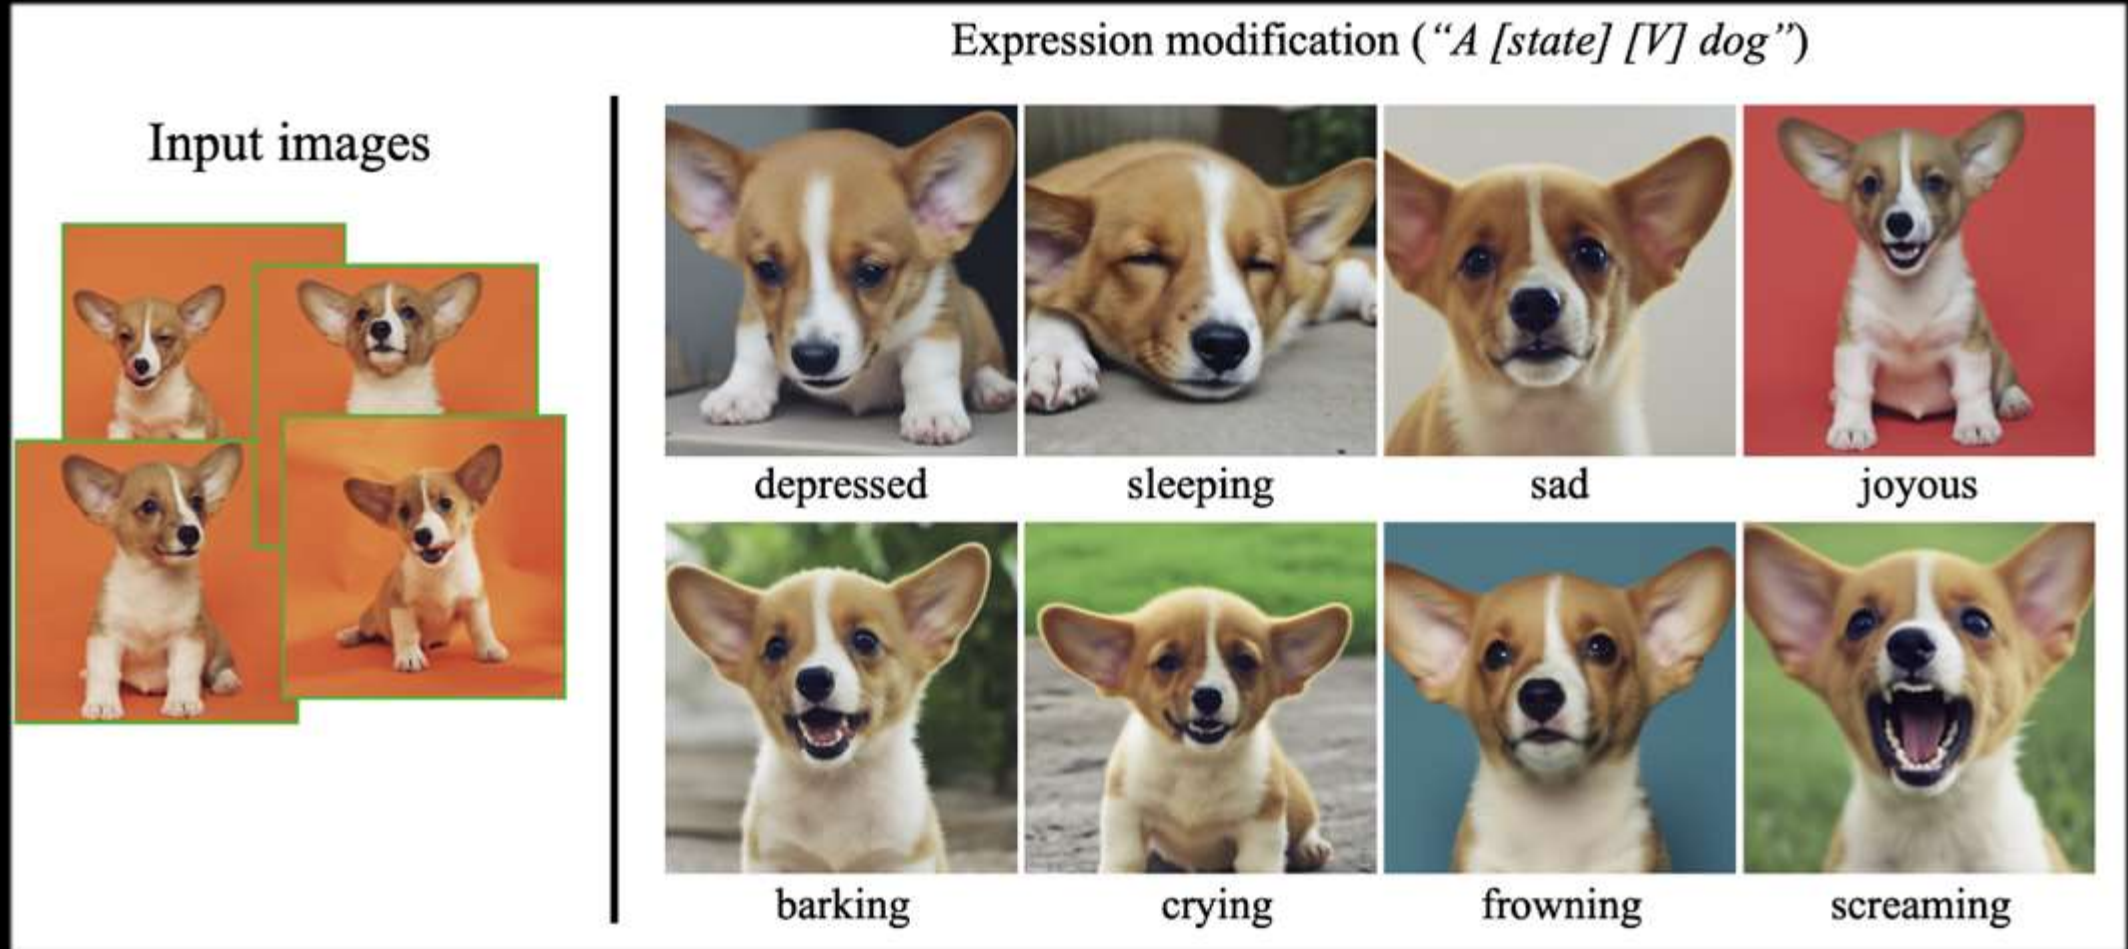
Pierre-Auguste Renoir



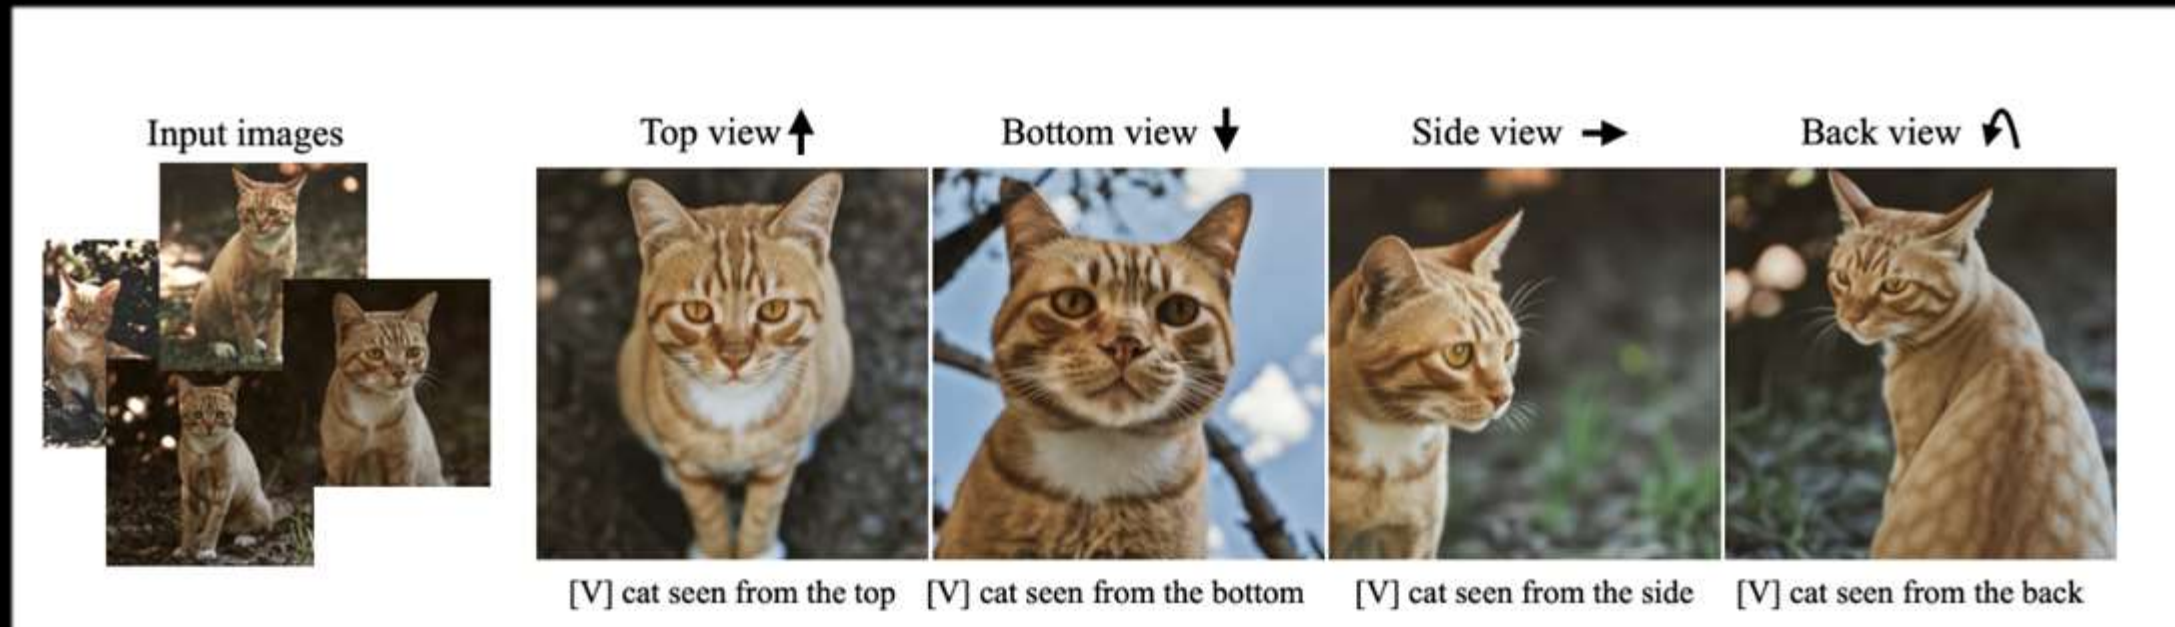
Leonardo da Vinci



# Results: Expression Manipulation



# Results: View Synthesis



# Results: Accessorization

Input images



Chef Outfit



Witch Outfit



Ironman Outfit



Nurse Outfit



Purple Wizard Outfit



Superman Outfit



Police Outfit



Angel Wings



# Results: Property Modification

Color modification (“A [color] [V] car”)



Input



purple



red



yellow



blue



pink

Hybrids (“A cross of a [V] dog and a [target species]”)



Input



bear



panda



koala



lion



hippo

# Main Failures

Input images



(a) Incorrect context synthesis



in the ISS



on the moon

(b) Context-appearance entanglement



in the Bolivian salt flats



on top of a blue fabric

(c) Overfitting



in the forest

# Reference

- *DreamBooth: Fine Tuning Text-to-Image Diffusion Models for Subject-Driven Generation*
- *Generative Adversarial Nets*
- *Denoising Diffusion Probabilistic Models*
- *High-Resolution Image Synthesis With Latent Diffusion Models*
- *An Image is Worth One Word: Personalizing Text-to-Image Generation using Textual Inversion*
- <https://dreambooth.github.io/>
- <https://jalammar.github.io/illustrated-stable-diffusion/>
- <https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>
- <https://medium.com/@steinsfu/stable-diffusion-clearly-explained-ed008044e07e>
- <https://theaisummer.com/diffusion-models/>
- <https://medium.com/@kemalpiro/step-by-step-visual-introduction-to-diffusion-models-235942d2f15c>
- <https://kailashahirwar.medium.com/a-very-short-introduction-to-diffusion-models-a84235e4e9ae>
- <https://erdem.pl/2023/11/step-by-step-visual-introduction-to-diffusion-models#forward-diffusion-diagram>



# Acknowledgement

***Thank you for your time and listening!***